

CSC411 Fall 2017

Assignment 2 Report

Tianbao Li

2017/11/04

1 Q1: Class-conditional gaussians

Given:

$$p(y = k) = a_k \quad (1)$$

$$p(\mathbf{x}|y = k, \mu, \sigma) = \left(\prod_{i=1}^d 2\pi\sigma_i^2\right)^{-1/2} \exp\left\{-\sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i - \mu_{ki})^2\right\} \quad (2)$$

1.1 Bayes' rule derivation

$$\begin{aligned} p(y = k|\mathbf{x}, \mu, \sigma) &= \frac{p(\mathbf{x}|y = k, \mu, \sigma)p(y = k)}{p(\mathbf{x}|\mu, \sigma)} \\ &= \frac{p(\mathbf{x}|y = k, \mu, \sigma)p(y = k)}{\sum_{j=1}^K p(\mathbf{x}|y = j, \mu, \sigma)} \\ &= \frac{(\prod_{i=1}^d 2\pi\sigma_i^2)^{-1/2} \exp\{-\sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i - \mu_{ki})^2\} a_k}{\sum_{j=1}^K (\prod_{i=1}^d 2\pi\sigma_i^2)^{-1/2} \exp\{-\sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i - \mu_{ji})^2\}} \end{aligned}$$

1.2 Negative likelihood function (NLL)

$$\begin{aligned} \ell(\theta; D) &= -\log p(y^{(1)}, \mathbf{x}^{(1)}, y^{(2)}, \mathbf{x}^{(2)}, \dots, y^{(N)}, \mathbf{x}^{(N)}|\theta) \\ &= -\log \prod_{n=1}^N p(y^{(n)}, x^{(n)}|\theta) \\ &= -\sum_{n=1}^N (\log p(x^{(n)}|y^{(n)}, \theta) + \log p(y^{(n)}|\theta)) \\ &= -\sum_{n=1}^N (\log((\prod_{i=1}^d 2\pi\sigma_i^2)^{-1/2} \exp\{-\sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i^{(n)} - \mu_{ki})^2\}) + \log \alpha_k) \\ &= -\sum_{n=1}^N (-\frac{1}{2} \sum_{i=1}^d 2\pi\sigma_i^2 - \sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i^{(n)} - \mu_{ki})^2 + \log \alpha_k) \\ &= \sum_{n=1}^N (\sum_{i=1}^d \pi\sigma_i^2 + \sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i^{(n)} - \mu_{ki})^2 - \log \alpha_k) \end{aligned}$$

1.3 Partial derivatives

$$\begin{aligned}
\frac{\partial \ell}{\partial \mu_{ki}} &= \frac{\partial (\sum_{n=1}^N \sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i^{(n)} - \mu_{ki})^2)}{\partial \mu_{ki}} \\
&= - \sum_{n=1}^N \sum_{i=1}^d \frac{1}{\sigma_i^2} \mathbb{1}(y^{(n)} = k) (x_i^{(n)} - \mu_{ki}) \\
\frac{\partial \ell}{\partial \sigma_i^2} &= \frac{\partial (\sum_{n=1}^N (\sum_{i=1}^d \pi \sigma_i^2 + \sum_{i=1}^d \frac{1}{2\sigma_i^2} (x_i^{(n)} - \mu_{ki})^2))}{\partial \sigma_i^2} \\
&= \sum_{n=1}^N \sum_{i=1}^d \mathbb{1}(y^{(n)} = k) (2\pi - \frac{1}{2\sigma_i^4} (x_i^{(n)} - \mu_{ki})^2)
\end{aligned}$$

1.4 Estimation

$$\begin{aligned}
\frac{\partial \ell}{\partial \mu_{ki}} &= 0 \\
\Rightarrow \mu_{ki} &= \frac{\sum_{n=1}^N \sum_{i=1}^d \mathbb{1}(y^{(n)} = k) x_i^{(n)}}{\sum_{n=1}^N \sum_{i=1}^d \mathbb{1}(y^{(n)} = k)} \\
\frac{\partial \ell}{\partial \sigma_i^2} &= 0 \\
\Rightarrow \sigma_i &= \sqrt[4]{\frac{\sum_{n=1}^N \sum_{i=1}^d \mathbb{1}(y^{(n)} = k) * (x_i^{(n)} - \mu_{ki})^2}{\sum_{n=1}^N \sum_{i=1}^d \mathbb{1}(y^{(n)} = k) * 4\pi}}
\end{aligned}$$

2 Handwritten digit classification

In this assignment, we use the dataset of handwritten digit. To get a macro view of the data, here we plot the means for each data classes in the training set, shown in Figure 1.

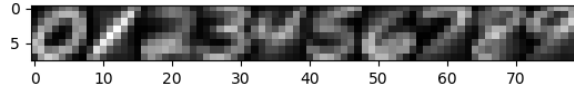


Figure 1: Mean for each class

2.1 K-NN classifier

First, we build a K nearest neighbor classifier using Euclidean distance on the data.

2.1.1 Accuracy for certain K values

- for K = 1
 - train classification accuracy: 1.0
 - test classification accuracy: 0.96875
- for K = 15

- train classification accuracy: 0.961
- test classification accuracy: 0.959

2.1.2 Ties solution

During the implementation, one important problem is that ties usually happens and need to be broken. The solution is shown as follows.

```
def query_knn(self, test_point, k):
    digit = None
    dist = self.l2_distance(test_point)
    noTies = False
    while noTies == False:
        mink = np.argpartition(dist, k)[:k]
        counts = np.bincount(self.train_labels[mink].astype(int))
        digit = np.argmax(counts == np.amax(counts))
        if len(digit) > 1:
            k = k + 1
        else:
            noTies = True
    return digit
```

While running KNN, if ties happen (which means two classes has the same maximum count), we try to decrease k by 1 and run the KNN again until no ties. As the class label should focus on one or a few classes, ties are unlikely to happen when $K = 1$, so this should be a solution to get an answer without ties.

2.1.3 Best k

Here we implement 10 fold cross validation in range 1-15 to find the optimal K. By finding minimum validation error, optimal K for this problem is that $K = 3$ and accuracies are shown as follows.

- train classification accuracy: 1.0
- validation classification accuracy: 0.96514286
- test classification accuracy: 0.96875

2.2 Conditional gaussian classifier training

2.2.1 Diagonal covariance

The diagonal elements of each covariance matrix Σ_k can be plotted as Figure 2.

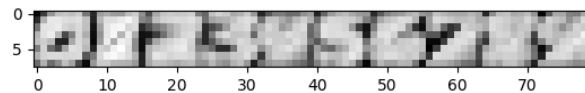


Figure 2: Diagonal elements of covariance for each class

2.2.2 Average conditional log-likelihood

The average conditional log-likelihood for each class is shown as follows.

- training data avg conditional likelihood: [-2.78902836 -2.48009327 -2.46864922 -2.53467722 -2.49894322 -2.93938866 -2.4646036 -3.70354979 -2.47281232 -2.74221317]
- test data avg conditional likelihood: [-3.22599621 -4.13098548 -2.44591828 -3.1675836 -2.65604537 -3.69867522 -3.67902811 -7.62636215 -2.81381993 -3.95651773]

2.2.3 Accuracy

By selecting the most likely posterior class for each datapoint as prediction, the accuracy is shown as follows.

- training accuracy: 0.981285714286
- test accuracy: 0.9605

2.2.4 Leading eigenvectors

The leading eigenvectors of each covariance matrix Σ_k can be plotted as Figure 3.

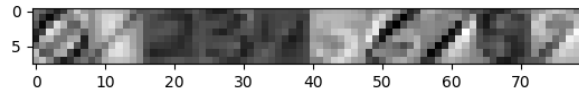


Figure 3: Leading eigenvectors of covariance for each class

2.3 Naive bayes classifier training

2.3.1 Convert to binart features

```
def binarize_data(pixel_values):  
    return np.where(pixel_values > 0.5, 1.0, 0.0)
```

2.3.2 Model fitting

For a Bernoulli Naive Bayes classifier using MAP, we need to calculate η first.

```
def compute_parameters(train_data, train_labels):  
    eta = np.zeros((10, 64))  
    K = eta.shape[0]  
    d = eta.shape[1]  
    for k in range(K):  
        k_index = np.argwhere(train_labels == k)  
        k_digits = train_data[k_index].reshape(-1, d)  
        for j in range(d):  
            eta[k][j] = (np.sum(k_digits[:, j]) + 1.0) \\  
                        / (k_digits.shape[0] + 2.0)  
    return eta
```

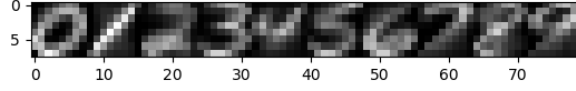


Figure 4: η for each class

2.3.3 Plot η

The parameters η is plotted in Figure 4.

2.3.4 Sample new data

Given the parameters η , we can sample new data points for each of the 10 digits. *numpy.random.binomial* method is used here. One new data point is like Figure 5.

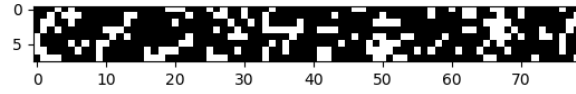


Figure 5: Sampled data point

2.3.5 Average conditional log-likelihood

The average conditional log-likelihood for each class is shown as follows.

- training data avg conditional likelihood: [-2.8909254 -3.92355812 -3.26626627 -3.10729808 -3.01401521 -3.12587265 -2.94008982 -3.17888068 -3.47393427 -3.54254905]
- test data avg conditional likelihood: [-3.10584429 -3.6500495 -3.3475608 -3.29021075 -3.04145573 -3.20550698 -3.06688846 -3.23938846 -3.44419131 -3.507459]

2.3.6 Accuracy

By selecting the most likely posterior class for each datapoint as prediction, the accuracy is shown as follows.

- training accuracy: 0.774142857143
- test accuracy: 0.76425